

Xiaoyang Cao

Cambridge, MA · (+1) 617-216-8184 · xycao@mit.edu · Personal Website · Google Scholar · GitHub

EDUCATION

Massachusetts Institute of Technology

M.S. in Computational Science & Engineering

Concurrent M.S. in Technology & Policy

Laboratory for Information and Decision Systems (LIDS)

Research advisor: Prof. [Michiel Bakker](#)

Cambridge, MA

Sep 2025 – May 2027 (expected)

Tsinghua University

B.S. in Mathematics and Physics

Beijing, China

Sep 2021 – Jun 2025

University of California, Berkeley

Exchange Student

Berkeley, CA

Jan 2024 – Aug 2024

RESEARCH INTERESTS

How intelligence emerges in collectives of interacting agents (multi-agent reinforcement learning, populations of LLM agents) and in self-referential learning (self-play, self-distillation, preference optimization).

PUBLICATIONS

- [1] **Xiaoyang Cao**, Jingqi Li, Zhe Fu, Alexandre M. Bayen. “Who Bears the Burden? Learning Responsibility for Shared Constraints in Multi-Agent Reinforcement Learning.” *Under review*, 2026. [arXiv:2610.07491](#).
- [2] **Xiaoyang Cao**, Siddarth Srinivasan, Michiel A. Bakker. “Do Modules Stay in Their Lane? Role Drift in Compound LLM Systems.” *Under review*, 2026. [arXiv:2607.21627](#).
- [3] **Xiaoyang Cao**, Zelai Xu, Mo Guang, Kaiwen Long, Michiel Bakker, Yu Wang, Chao Yu. “RE-PO: Robust Enhanced Policy Optimization as a General Framework for LLM Alignment.” *International Conference on Learning Representations (ICLR)*, 2026.
- [4] **Xiaoyang Cao**, Zhe Fu, Alexandre M. Bayen. “Pareto Control Barrier Function for Inner Safe Set Maximization Under Input Constraints.” *American Control Conference (ACC)*, 2025.
- [5] **Xiaoyang Cao**, Dingyi Zhuang, Shenhao Wang, Jinhua Zhao. “Virtual Nodes Improve Long-term Traffic Prediction.” *Transportation Research Board Annual Meeting (TRB)*, 2025.

RESEARCH EXPERIENCE

Social Swarm: Scaling Populations of Research Agents with Sparse Networks

Sep 2026 – Present

Advised by Prof. [Michiel Bakker](#), MIT

- Studying how collectives of LLM research agents convert agent count into capability on open-ended problems, motivated by recent results that scale to thousands of agents.
- Identified *diversity collapse*: when every agent sees every attempt, the population homogenizes and long-horizon performance stalls. Proposed a sparse follow-network organization in which each agent sees only a bounded set of peers.
- Preliminary experiments on the Heilbronn triangle problem show that the sparse network delays collapse relative to a fully visible shared board and to independent best-of- N .

Learning Responsibility for Shared Constraints in MARL [1] [[arXiv](#)] [[Project](#)]

May 2025 – Sep 2026

Advised by Prof. [Alexandre Bayen](#), UC Berkeley

- Proposed LiRA (Lagrangian Responsibility Allocation): a common multiplier enforces a shared cost budget, while learned per-agent responsibility shares redistribute its penalty by optimizing social welfare, without modifying rewards or constraints.
- Proved that for convex games varying the shares induces a smooth family of normalized generalized Nash equilibria

in which active constraints stay at budget while welfare varies, and derived a welfare gradient that accounts for both learning updates and the induced shift in data distribution.

- Improved average social welfare by up to 29% over uniform and agent-specific multiplier baselines across CityLearn, MABIM, Harvest, and MetaDrive (3 to 400 agents).

Role Drift in Compound LLM Systems [2] [arXiv] [VentureBeat]

Apr 2026 – Sep 2026

Advised by Prof. [Michiel Bakker](#), MIT

- Identified *role drift*: end-to-end reinforcement learning can improve terminal performance while causing specialized modules to abandon their intended functions.
- Developed role-level evaluations showing that 86% of an apparent 31-percentage-point accuracy gain disappeared once the decomposer’s assigned function was measured, demonstrating that end-to-end metrics can conceal internal failure.
- Proposed Role Anchor, a training objective that preserves the influence of role specifications while allowing modules to improve under end-to-end optimization.

RE-PO: Robust LLM Preference Alignment (ICLR 2026) [3] [arXiv] [Project]

Mar 2025 – Sep 2025

Advised by Prof. [Yu Wang](#), Tsinghua University

- Proposed an EM-based framework that models preference-label correctness as a latent variable, jointly estimates label and annotator reliability, and systematically robustifies DPO, IPO, SimPO, and CPO.
- Proved convergence to the true noise level under calibration and improved AlpacaEval 2 win rates by up to 7.0 percentage points on Mistral-7B and Llama-3-8B, with additional gains on Arena-Hard and MultiPref.

Safe Control via Pareto Control Barrier Functions (ACC 2025) [4] [arXiv]

May 2024 – Oct 2024

Advised by Prof. [Alexandre Bayen](#), UC Berkeley

- Developed the Pareto Control Barrier Function algorithm by integrating Pareto multi-task learning into neural control barrier functions, maximizing the inner safe set while guaranteeing safety under input constraints.
- Matched the Hamilton–Jacobi viability kernel on a 2D inverted pendulum and achieved a 700% larger inner safe set than the benchmark on a 12-D quadrotor obstacle-avoidance task.

Virtual Nodes for Long-term Traffic Prediction (TRB 2025) [5] [arXiv]

May 2024 – Sep 2024

Advised by Prof. [Jinhua Zhao](#), MIT

- Introduced virtual nodes into spatio-temporal graph neural networks to mitigate over-squashing in long-horizon traffic forecasting, using a semi-adaptive adjacency matrix learned from node embeddings.
- Reduced RMSE by 6.27% and MAPE by 5.04% on a five-year, 716-sensor San Diego dataset; the learned adjacency emphasized traffic-intensive intersections, improving interpretability.

AWARDS & HONORS

Outstanding Winner & INFORMS Award, Mathematical Contest in Modeling – top 1%

May 2023

Outstanding Innovation in Science & Technology Scholarship, Tsinghua University

2023, 2024

Comprehensive Excellence Scholarship, Tsinghua University

Oct 2023

Outstanding Academic Scholarship, Tsinghua University

Oct 2022

SERVICE & TEACHING

Conference Reviewer: ICLR 2027, CoRL 2026

Teaching Assistant: Calculus A, Tsinghua University

Spring 2025

SKILLS & LANGUAGES

ML & RL: Python, PyTorch, Hugging Face Transformers / TRL, vLLM, DeepSpeed, FSDP, PettingZoo, RLlib

Compute: Multi-GPU, multi-node training on H100 / H200 / B200 clusters

Languages: English (fluent), Chinese (native)